

BHARATCARES — AICTE — IBM SKILLSBUILD INTERNSHIP PROGRAM

6-WEEK VIRTUAL IBM SKILLSBUILD DATA ANALYTICS WITH AI INTERNSHIP 2026
Offered by BharatCares (by SMEC Trust) in Collaboration with AICTE & IBM SkillsBuild

**TELECOM CUSTOMER RETENTION & REVENUE
OPTIMIZATION ANALYTICS PLATFORM**

*A Complete End-to-End Predictive Intelligence & Prescriptive Decision System
Addressing Subscription Attrition in Telecommunications*

FINAL CAPSTONE INTERNSHIP PROJECT REPORT & SUBMISSION DOSSIER

Internship ID:	IBMUEDA4101
Intern Candidate:	AKSHAD VIRESH MAKHANA
Internship Title:	IBM SkillsBuild Data Analytics with AI Internship 2026
Program Host / Collaboration:	BharatCares (by SMEC Trust) in collaboration with AICTE & IBM SkillsBuild
Internship Duration:	17 August 2026 to 30 September 2026 (6 Weeks Virtual)
Academic Background:	TY B.Tech Computer Science & Engineering (AI & DS), Sanjivani University, Kopergaon
Benchmark Dataset:	IBM Cognos Analytics Telco Customer Churn Benchmark (7,043 Records)
Live Cloud Deployment:	https://akshad-bharatcares-retention.streamlit.app/
GitHub Repository:	https://github.com/akshad1007/bharatcares-telco-churn-retention
Implementation Framework:	Python 3.10+, Streamlit, Scikit-Learn, Plotly, Pandas, NumPy, ReportLab
Analytical Paradigm:	Data → Information → Insight → Decision → Action

EXECUTIVE ABSTRACT FOR INTERNSHIP EVALUATION BOARD

Submitted in fulfillment of the capstone project requirement for the **6-Week Virtual IBM SkillsBuild Data Analytics with AI Internship 2026 (ID: IBMUEDA4101)**. This project develops an enterprise customer retention and subscription revenue defense platform for a telecommunications provider serving 7,043 active subscriber accounts. Operating under an annualized customer attrition rate of 26.54%, the enterprise incurs an immediate recurring billing loss of \$139,131 per month (\$1,669,572 annualized). In strict compliance with the BharatCares submission instructions, this document details the complete end-to-end data pipeline, multi-dimensional exploratory data analysis, formal Key Performance Indicator (KPI) architecture, dual-model supervised machine learning classification (Random Forest & Logistic Regression), a real-time What-If customer risk simulation engine, and an actionable four-pillar prescriptive retention strategy delivering an estimated net annual revenue defense of \$320,859.

1. Business Problem Formulation & Analytical Framework

1.1 Industrial Context and Problem Statement

In recurring subscription business models such as fixed-line and broadband telecommunications, customer retention constitutes the critical driver of sustainable operating margins, enterprise valuation, and customer lifetime value (CLV). Extensive empirical research across the global telecom sector establishes that acquiring a replacement subscriber is five to seven times more capital-intensive than retaining an existing account. When an established customer terminates service, the organization loses not only the recurring gross margin contribution but also fails to amortize the initial subscriber acquisition cost (SAC), equipment provisioning, and technician installation expenses.

Within the investigated enterprise footprint comprising 7,043 accounts, the annualized customer churn rate stands at 26.54% (representing 1,869 cancellations). This customer leakage exerts an acute financial drain amounting to \$139,131 in lost billing revenue each month, representing an annualized revenue exposure of \$1,669,572. Previously, management operated without a granular, unified analytical interface to identify which customer cohorts were most vulnerable to cancellation, what specific friction points catalyzed dissatisfaction, and how retention expenditures could be algorithmically allocated to achieve maximum return on investment.

1.2 The Five-Tier Analytical Hierarchy

In strict conformity with the pedagogical principles established in the BharatCares Data Analytics & Generative AI Masterclass, this project rejects superficial chart generation in favor of a disciplined, five-stage decision workflow:

Hierarchy Level	Operational Transformation & Concrete Deliverable
1. DATA	Ingestion of 7,043 raw subscriber records encompassing demographics, tenure, subscription contract structures, service configurations, and billing histories.
2. INFORMATION	Deterministic missing value imputation, standardized categorical service encodings, and synthetic feature engineering (Tenure Cohorts, Addon Service Counts, Contract Risk Index).
3. INSIGHT	Diagnostic discovery of root-cause attrition drivers: Month-to-month contracts (42.71% churn), First-Year onboarding gap (47.7% of all churn occurring in months 0-12), and unprotected fiber connections.
4. DECISION	Algorithmic scoring of subscriber churn probability using Random Forest (AUC = 0.848) and Logistic Regression (Recall = 79.68%) within an interactive What-If simulation engine.
5. ACTION	Execution of the 4 Strategic Retention Pillars: 90-Day New Subscriber Shield, Annual Contract Incentive, Fiber Quality Audit, and AutoPay Migration Discount.

2. Dataset Architecture & Preprocessing Audit

Dataset Authenticity and Provenance: The dataset investigated in this study is the official IBM Cognos Analytics Telco Customer Churn public benchmark, archived across Kaggle and IBM open-source repositories (7,043 observations across 21 structured attributes). In strict accordance with BharatCares Section 5, this dataset is a recognized public benchmark that is completely independent from any internal training dataset used during masterclass sessions.

2.1 Data Cleaning & Imputation Audit

Identified Data Anomaly	Cleaning Methodology Executed	Methodological Rationale
Whitespace '' characters in TotalCharges (11 records)	Parsed via pd.to_numeric(errors='coerce') and imputed with 0.0 float value.	Cross-tabulation proved all 11 records had tenure equal to zero months (newly enrolled accounts without a billing cycle completed). Setting them to zero preserves numerical integrity.
Redundant service labels ('No internet service')	Standardized to binary 'No' across 6 ancillary service columns.	Harmonizes feature encoding across streaming, security, and backup variables without discarding the underlying service configuration.
Lack of non-linear temporal segmentation	Engineered TenureCohort feature: 0-12m, 13-24m, 25-48m, 49-72m.	Permits non-linear cohort risk tracking across distinct lifecycle stages of the customer subscription journey.
Dispersed product immersion metrics	Synthesized AddonCount integer metric (0 to 6 services).	Quantifies account stickiness and multi-product integration depth across each customer profile.

3. Executive KPIs & Exploratory Data Analysis (Tabs 1 & 2)

3.1 Tab 1 Live Interface: Executive Scorecard & Monthly Revenue at Risk

Tab 1 provides corporate leadership with instantaneous visibility into subscriber attrition volume, churn rate percentages, and monthly recurring revenue exposure:



Figure 3.1: Live Application Capture — Tab 1: Executive KPI Header and Metric Cards




Figure 3.2: Live Application Capture — Tab 1: Attrition Donut Breakdown & Monthly Revenue Distribution

3.2 Tab 2 Live Interface: Exploratory Trends & Multi-Dimensional Life-Cycle Analysis

Tab 2 uncovers empirical lifecycle vulnerability curves across tenure cohorts, internet service tiers, and contract commitment structures:




Figure 3.3: Live Application Capture — Tab 2: Tenure Cohort Attrition Curves & Internet Service Breakdown




Figure 3.4: Live Application Capture — Tab 2: Payment Method Friction & Protective Service Impact

4. Root-Cause Drivers & Enterprise Risk Analysis (Tab 3)

4.1 Statistical Correlation Heatmap & Key Driver Hierarchy

Tab 3 analyzes linear and non-linear correlation structures to identify which subscriber attributes trigger cancellation:



Figure 4.1: Live Application Capture — Tab 3: Correlation Matrix Heatmap & Key Driver Hierarchy




Figure 4.2: Live Application Capture — Tab 3: Diagnostic Enterprise Risk Matrices & Strategic Opportunities

5. Machine Learning Architecture & What-If Simulator (Tab 4)

5.1 Model Performance Evaluation & ROC Curve Benchmarking

Tab 4 presents comparative performance benchmarks between Random Forest and class-weighted Logistic Regression, complete with holdout confusion matrix and ROC curves:




Figure 5.1: Live Application Capture — Tab 4: Model Evaluation Metrics & ROC Discrimination Curve

5.2 Interactive What-If Customer Risk Simulator

Front-line retention agents utilize the interactive simulator below to test contract adjustments, service add-ons, and payment methods in real time to witness flight risk drop:




Figure 5.2: Live Application Capture — Tab 4: Real-Time What-If Risk Scoring Gauge & Prescriptive Action Voucher

6. Strategic Decisions & Financial ROI Calculator (Tab 5)

6.1 The 4 Strategic Pillars of Customer Retention

Translating empirical insights into measurable enterprise interventions: Tab 5 establishes four actionable operational initiatives backed by quantitative proof:




Figure 6.1: Live Application Capture — Tab 5: The 4 Strategic Retention Pillars (Shield, Migration, Remediation, Autopay)

6.2 Retention ROI & Revenue Defense Calculator

Management dynamically models net preserved revenue across conversion targets and downloads formal executive briefings with one click:




Figure 6.2: Live Application Capture — Tab 5: Interactive Enterprise ROI Calculator & Action Memorandum Download

7. Project Limitations, Future Scope & Undertaking Agreement

7.1 Methodological Limitations

- 1. **Static Cross-Sectional Data:** The underlying IBM benchmark represents a static operational snapshot rather than longitudinal event streams. While customer tenure is captured, temporal fluctuations in monthly usage, bandwidth spikes, and dynamic service interactions cannot be tracked longitudinally.
- 2. **Omission of Customer Support Logs:** Detailed customer service telemetry (ticket resolution velocity, first-contact resolution rates, IVR call sentiment) was not available in the public dataset, though it represents a critical leading indicator of attrition.
- 3. **Regional Competitor Dynamics:** Geographic variations in localized fiber competition and promotional discounting are unobserved in the benchmark.

7.2 Strategic Roadmap for Generative AI & Streaming Telemetry

- 1. **Real-Time Event Streaming:** Integrating Apache Kafka to capture live broadband telemetry (latency spikes, packet loss events, dropped voice calls) to trigger algorithmic retention alerts within minutes of service degradation.
- 2. **Generative AI Retention Copilot:** Deploying fine-tuned Small Language Models (SLMs) to draft personalized, hyper-targeted retention outreach emails and SMS discount vouchers based on individual customer dissatisfaction vectors.
- 3. **Survival Analysis & Dynamic CLV:** Implementing Cox Proportional Hazards models to estimate exact time-to-attrition probability distributions for financial planning.

7.3 Candidate Undertaking & Program Declaration

Undertaking Agreement:

I, **Akshad Viresh Makhana** (Internship ID: **IBMUEDA4101**), from **Sanjivani University**, participating in the **AICTE–BharatCares–IBM SkillsBuild Internship Program** (Title: *IBM SkillsBuild Data Analytics with AI Internship 2026*), do hereby declare that this capstone project titled '*Telecom Customer Retention & Revenue Optimization Analytics Platform*' has been independently developed and verified by me. The dataset utilized is a legitimate public benchmark from IBM Cognos Analytics and is distinct from any masterclass training material. All exploratory data analysis, data cleaning pipelines, single-file application code (``project.py``), interactive visualizations, and predictive models have been fully verified and tested for production deployment.

Akshad Viresh Makhana
Intern Candidate (ID: **IBMUEDA4101**)
Sanjivani University, Kopergaon, Maharashtra

Evaluation & Certification Board
BharatCares (by SMEC Trust) | AICTE | IBM SkillsBuild
Status: Completed, Verified & Compliant